# DEVYANI BEOHAR

# J068

# Sklearn API for:

## LINEAR REGRESSION

*class* sklearn.linear_model.**LinearRegression**(**, fit_intercept=True, normalize=False, copy_X=True, n_jobs=None, positive=False*)

**fit_intercept***bool, default=True*
Whether to calculate the intercept for this model.

**normalize***bool, default=False*
This parameter is ignored when `fit_intercept` is set to False.

**copy_X***bool, default=True*
If True, X will be copied; else, it may be overwritten.

**n_jobs***int, default=None*
The number of jobs to use for the computation.

**positive***bool, default=False*
When set to `True`, forces the coefficients to be positive.

Ordinary least squares Linear Regression.

LinearRegression fits a linear model with coefficients w = (w1, ..., wp) to minimize the residual sum of squares between the observed targets in the dataset, and the targets predicted by the linear approximation.

# LOGISTIC REGRESSION

*class* `sklearn.linear_model.`**`LogisticRegression`**(*penalty='l2', *, dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1, class_weight=None, random_state=None, solver='lbfgs', max_iter=100, multi_class='auto', verbose=0, warm_start=False, n_jobs=None, l1_ratio=None*)

**penalty*{'l1', 'l2', 'elasticnet', 'none'}, default='l2'***
Used to specify the norm used in the penalization.

**dual*bool, default=False***
Dual or primal formulation.

**tol*float, default=1e-4***
Tolerance for stopping criteria.

**C*float, default=1.0***
Inverse of regularization strength; must be a positive float.

**fit_intercept*bool, default=True***
Specifies if a constant (a.k.a. bias or intercept) should be added to the decision function.

**intercept_scaling*float, default=1***
Useful only when the solver 'liblinear' is used and self.fit_intercept is set to True.

**class_weight*dict or 'balanced', default=None***
Weights associated with classes in the form `{class_label: weight}`. If not given, all classes are supposed to have weight one.

**random_state*int, RandomState instance, default=None***
Used when `solver` == 'sag', 'saga' or 'liblinear' to shuffle the data.

***solver='lbfgs'***
'lbfgs' handle multinomial loss

**max_iter*int, default=100***
Maximum number of iterations taken for the solvers to converge.

**multi_class*{'auto', 'ovr', 'multinomial'}, default='auto'***

If the option chosen is ‘ovr’, then a binary problem is fit for each label.

**verbose*int, default=0***
For the liblinear and lbfgs solvers set verbose to any positive number for verbosity.

**warm_start*bool, default=False***
When set to True, reuse the solution of the previous call to fit as initialization

**n_jobs*int, default=None***
Number of CPU cores used when parallelizing over classes if multi_class=’ovr’”.

**l1_ratio*float, default=None***
The Elastic-Net mixing parameter, with `0 <= l1_ratio <= 1`.

Logistic regression is a statistical analysis method used to predict a data value based on prior observations of a data set. Logistic regression can also play a role in data preparation activities by allowing data sets to be put into specifically predefined buckets during the extract, transform, load ETL process in order to stage the information for analysis.

# RIDGE

*class* `sklearn.linear_model.`**`Ridge`**(*alpha=1.0*, ***, *fit_intercept=True*, *normalize=False*, *copy_X=True*, *max_iter=None*, *tol=0.001*, *solver='auto'*, *random_state=None*)

**alpha*{float, ndarray of shape (n_targets,)}, default=1.0***
Regularization strength; must be a positive float.

**fit_intercept*bool, default=True***
Whether to fit the intercept for this model.

**normalize*bool, default=False***
This parameter is ignored when `fit_intercept` is set to False.

**copy_X*bool, default=True***
If True, X will be copied; else, it may be overwritten.

**max_iter*int, default=None***
Maximum number of iterations for conjugate gradient solver.

**tol*float, default=1e-3***
Precision of the solution.

**solver*{'auto'}***
'auto' chooses the solver automatically based on the type of data.

Ridge regression is a model tuning method that is used to analyse any data that suffers from multicollinearity. This method performs L2 regularization.
This model solves a regression model where the loss function is the linear least squares function and regularization is given by the l2-norm. Also known as Ridge Regression or Tikhonov regularization. This estimator has built-in support for multi-variate regression .

# LASSO

*class* `sklearn.linear_model.`**`Lasso`**(*alpha=1.0*, ***, *fit_intercept=True*, *normalize=False*, *precompute=False*, *copy_X=True*, *max_iter=1000*, *tol=0.0001*, *warm_start=False*, *positive=False*, *random_state=None*, *selection='cyclic'*)

**alpha*float, default=1.0***
Constant that multiplies the L1 term. Defaults to 1.0. `alpha = 0` is equivalent to an ordinary least square, solved by the **`LinearRegression`** object.

**fit_intercept*bool, default=True***
Whether to calculate the intercept for this model.

**normalize*bool, default=False***
This parameter is ignored when `fit_intercept` is set to False.

**precompute*bool or array-like of shape (n_features, n_features), default=False***
Whether to use a precomputed Gram matrix to speed up calculations.

**copy_X*bool, default=True***
If `True`, X will be copied; else, it may be overwritten.

**max_iter*int, default=1000***
The maximum number of iterations.

**ol*float, default=1e-4***

The tolerance for the optimization: if the updates are smaller than `to 1`, the optimization code checks the dual gap for optimality and continues until it is smaller than `to 1`

**warm_start*bool, default=False***

When set to True, reuse the solution of the previous call to fit as initialization, otherwise, just erase the previous solution.

**positive*bool, default=False***

When set to `True`, forces the coefficients to be positive.

**random_state*int, RandomState instance, default=None***

The seed of the pseudo random number generator that selects a random feature to update.

**selection*{'cyclic', 'random'}, default='cyclic'***

If set to 'random', a random coefficient is updated every iteration rather than looping over features sequentially by default.

Lasso regression is a regularization technique. It is used over regression methods for a more accurate prediction. This model uses shrinkage. Shrinkage is where data values are shrunk towards a central point as the mean.